Stage 1 Results

Dataset	Encoder	Method	K-shot	Runs	Test Accuracy
dtd	dinov2_vits14	linear_probe	5	3	63.01% +/- 0.27%
dtd	dinov2_vits14	linear_probe	10	3	68.58% +/- 0.80%
dtd	dinov2_vits14	linear_probe	full	3	76.35% +/- 0.72%
dtd	dinov2_vits14	prototype	5	3	64.33% +/- 0.30%
dtd	dinov2_vits14	prototype	10	3	68.46% +/- 1.01%
dtd	dinov2_vits14	prototype	full	1	72.29%
dtd	resnet18	linear_probe	5	3	45.80% +/- 0.75%
dtd	resnet18	linear_probe	10	3	52.39% +/- 0.74%
dtd	resnet18	linear_probe	full	3	62.55% +/- 0.51%
dtd	resnet18	prototype	5	3	46.93% +/- 0.32%
dtd	resnet18	prototype	10	3	51.97% +/- 0.54%
dtd	resnet18	prototype	full	1	58.78%
flowers102	resnet18	linear_probe	5	3	75.13% +/- 0.22%
flowers102	resnet18	linear_probe	10	3	83.22% +/- 0.14%
flowers102	resnet18	linear_probe	full	3	83.22% +/- 0.14%
flowers102	resnet18	prototype	5	3	69.62% +/- 0.32%
flowers102	resnet18	prototype	10	3	75.22% +/- 0.00%
flowers102	resnet18	prototype	full	1	75.22%

Accuracy vs. training-set size

dtd / dinov2_vits14

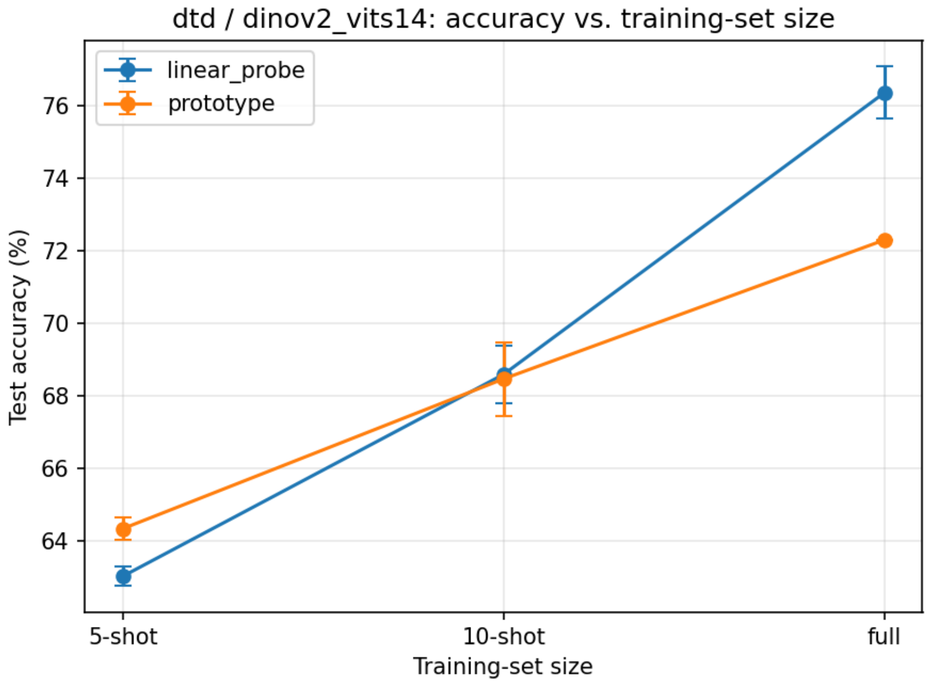

dtd / resnet18

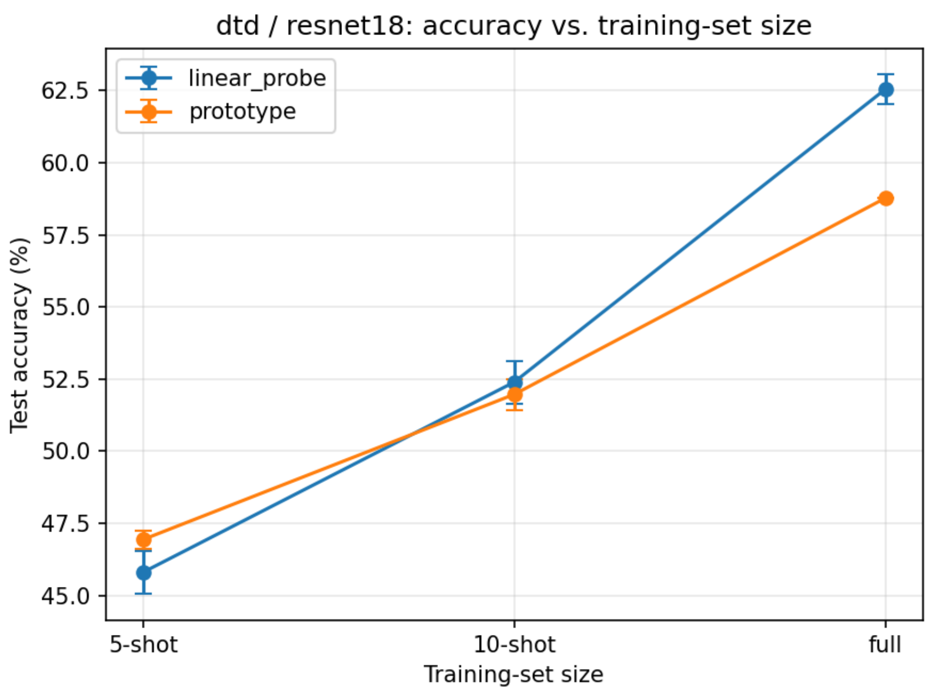

flowers102 / resnet18

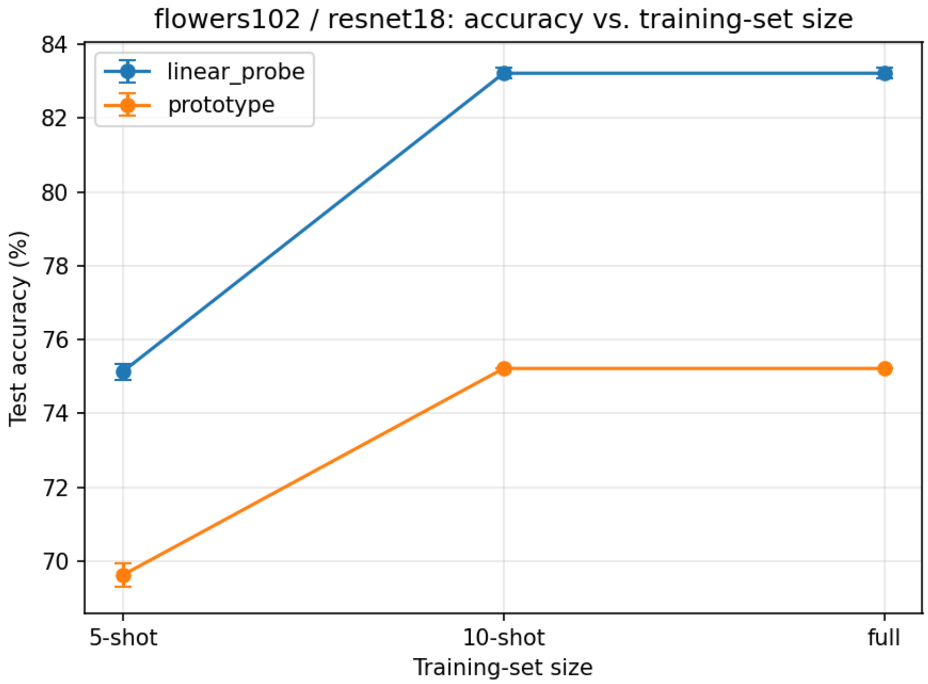

Training / validation loss (10-shot, seed 0)

dtd / dinov2_vits14

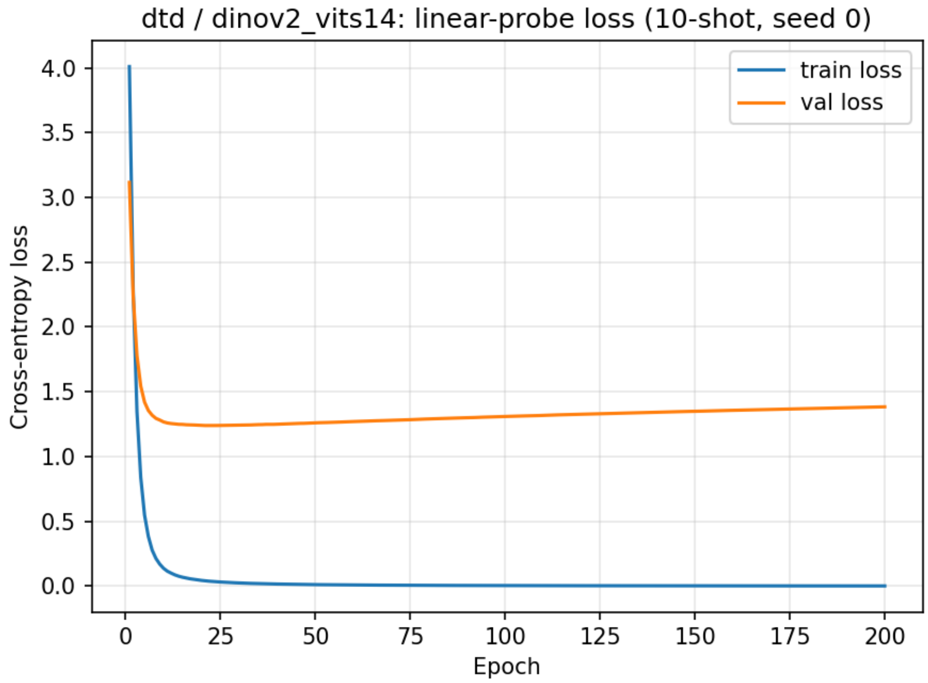

dtd / resnet18

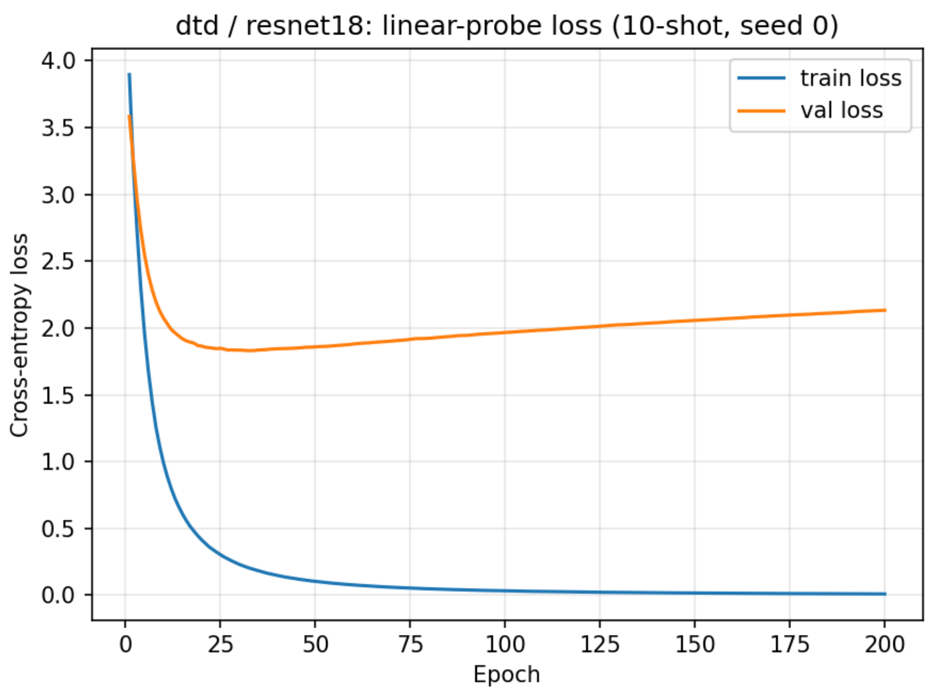

flowers102 / resnet18

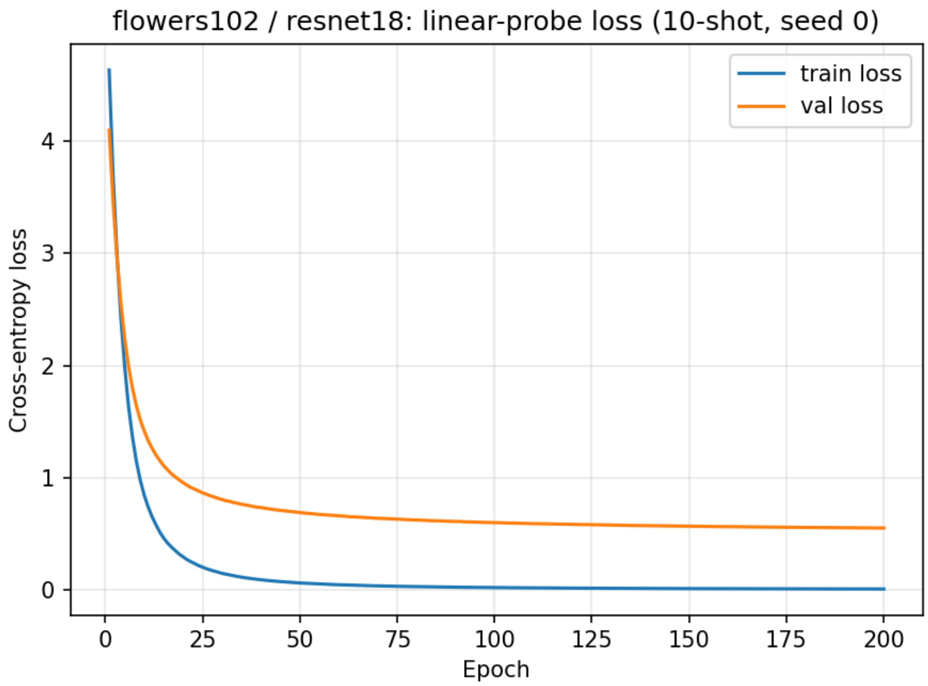

Confusion matrices (row-normalized)

dtd / dinov2_vits14

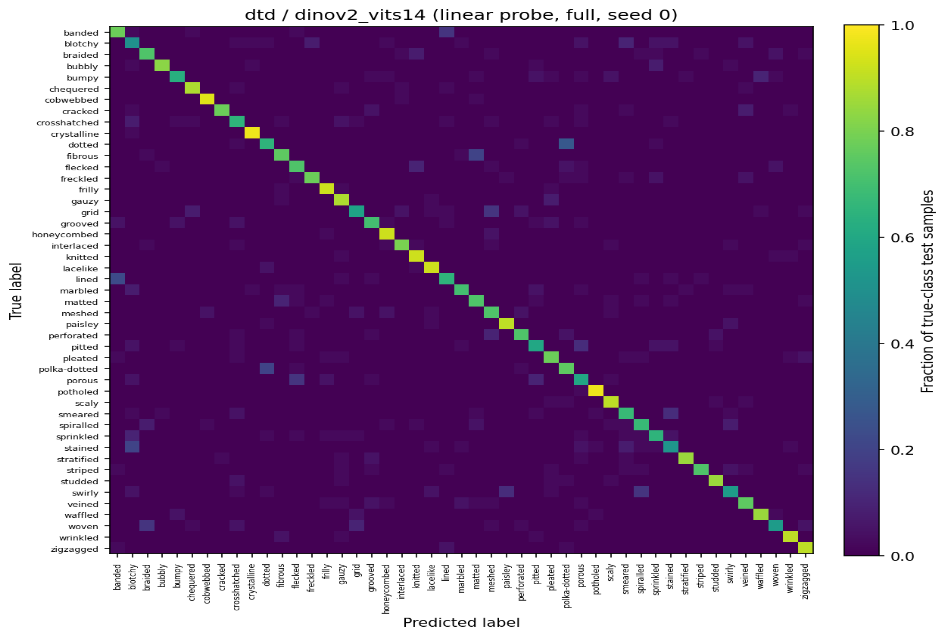

flowers102 / resnet18

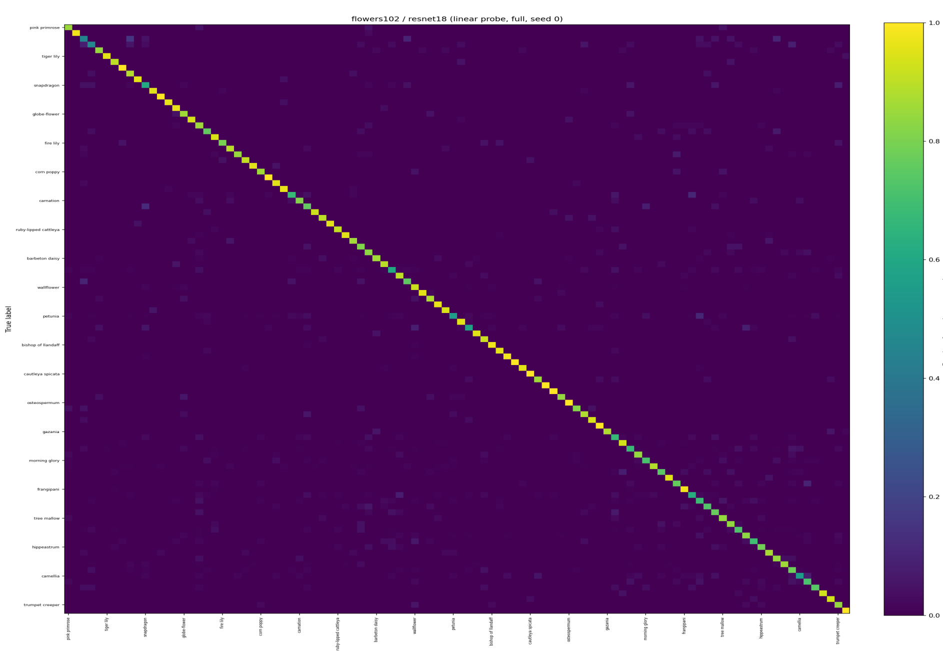

Feature-space visualizations (t-SNE)

dtd / dinov2_vits14

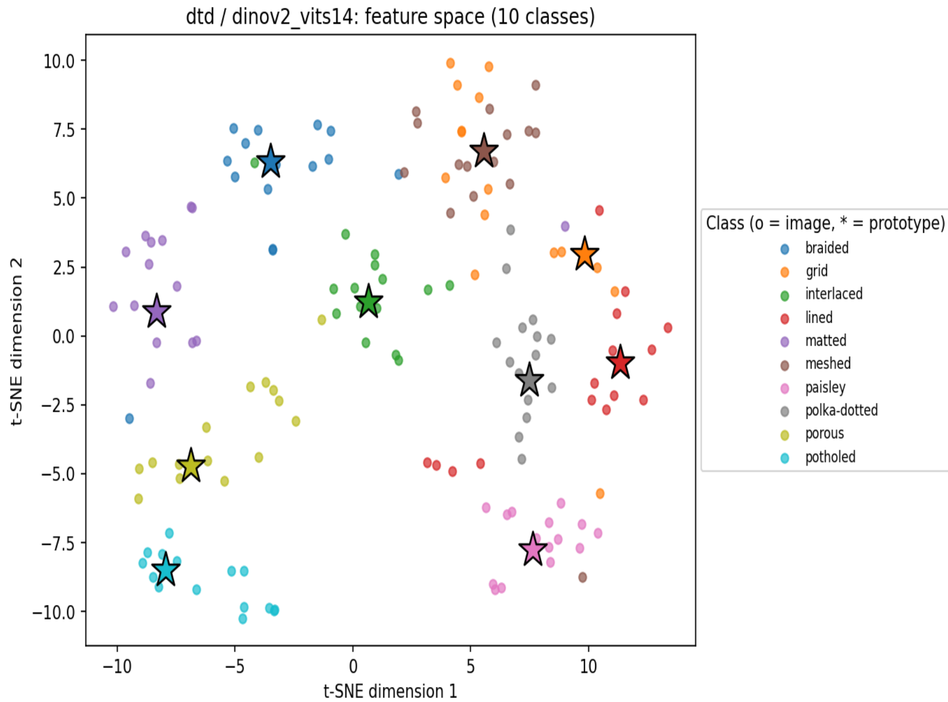

dtd / resnet18

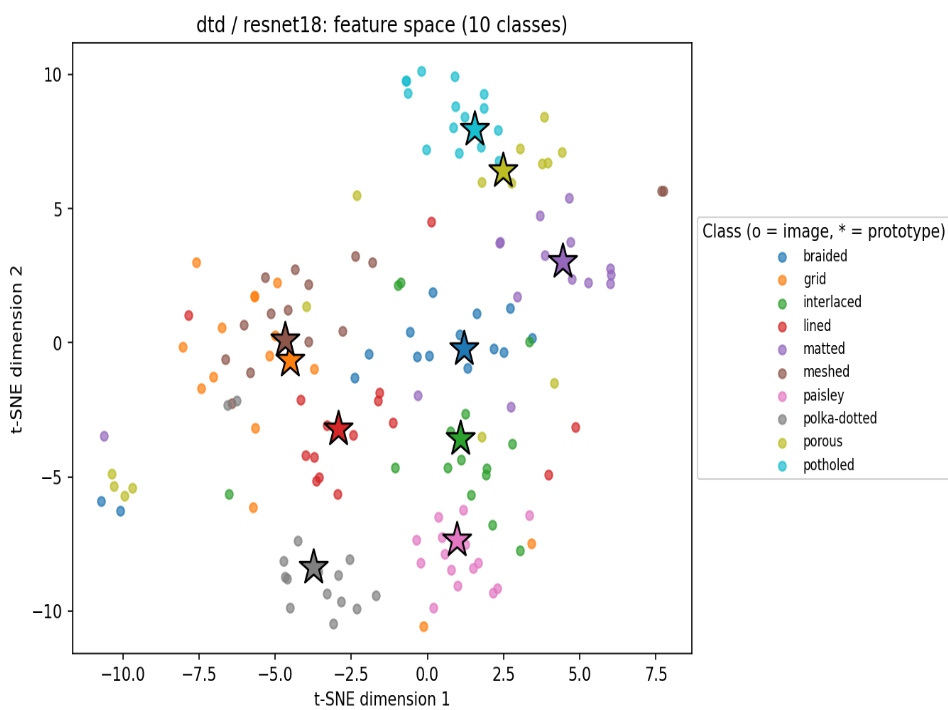

flowers102 / resnet18